

# Assignment\_1

August 21, 2024

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score, mean_squared_error

# Load the dataset
data = pd.read_csv("uber.csv")

data
```

```
[1]:
```

	Unnamed: 0	key	fare_amount	\
0	24238194	2015-05-07 19:52:06.0000003	7.5	
1	27835199	2009-07-17 20:04:56.0000002	7.7	
2	44984355	2009-08-24 21:45:00.00000061	12.9	
3	25894730	2009-06-26 08:22:21.0000001	5.3	
4	17610152	2014-08-28 17:47:00.000000188	16.0	
...	...	...	...	
199995	42598914	2012-10-28 10:49:00.00000053	3.0	
199996	16382965	2014-03-14 01:09:00.0000008	7.5	
199997	27804658	2009-06-29 00:42:00.00000078	30.9	
199998	20259894	2015-05-20 14:56:25.0000004	14.5	
199999	11951496	2010-05-15 04:08:00.00000076	14.1	

	pickup_datetime	pickup_longitude	pickup_latitude	\
0	2015-05-07 19:52:06 UTC	-73.999817	40.738354	
1	2009-07-17 20:04:56 UTC	-73.994355	40.728225	
2	2009-08-24 21:45:00 UTC	-74.005043	40.740770	
3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	
4	2014-08-28 17:47:00 UTC	-73.925023	40.744085	
...	...	...	...	
199995	2012-10-28 10:49:00 UTC	-73.987042	40.739367	
199996	2014-03-14 01:09:00 UTC	-73.984722	40.736837	
199997	2009-06-29 00:42:00 UTC	-73.986017	40.756487	

199998	2015-05-20 14:56:25 UTC	-73.997124	40.725452
199999	2010-05-15 04:08:00 UTC	-73.984395	40.720077

	dropoff_longitude	dropoff_latitude	passenger_count
0	-73.999512	40.723217	1
1	-73.994710	40.750325	1
2	-73.962565	40.772647	1
3	-73.965316	40.803349	3
4	-73.973082	40.761247	5
...	...	...	...
199995	-73.986525	40.740297	1
199996	-74.006672	40.739620	1
199997	-73.858957	40.692588	2
199998	-73.983215	40.695415	1
199999	-73.985508	40.768793	1

[200000 rows x 9 columns]

```
[2]: # 1. Pre-process the dataset

# Remove unnecessary column
data["pickup_datetime"] = pd.to_datetime(data["pickup_datetime"])

missing_values = data.isnull().sum()
print("Missing values in the dataset:")
print(missing_values)

# Handle missing values
# We can choose to drop rows with missing values or fill them with appropriate
# values.

data.dropna(inplace=True)

# To fill missing values with the mean value of the column:
# data.fillna(data.mean(), inplace=True)

# Ensure there are no more missing values
missing_values = data.isnull().sum()
print("Missing values after handling:")
print(missing_values)


# 2. Identify outliers
# visualization to detect outliers.
sns.boxplot(x=data["fare_amount"])
plt.show()
```

Missing values in the dataset:

```

Unnamed: 0      0
key             0
fare_amount     0
pickup_datetime 0
pickup_longitude 0
pickup_latitude 0
dropoff_longitude 1
dropoff_latitude 1
passenger_count 0
dtype: int64
Missing values after handling:
Unnamed: 0      0
key             0
fare_amount     0
pickup_datetime 0
pickup_longitude 0
pickup_latitude 0
dropoff_longitude 0
dropoff_latitude 0
passenger_count 0
dtype: int64

```



```

[3]: # Calculate the IQR for the 'fare_amount' column
Q1 = data["fare_amount"].quantile(0.25)
Q3 = data["fare_amount"].quantile(0.75)

```

```


IQR = Q3 - Q1

# Define a threshold (e.g., 1.5 times the IQR) to identify outliers
threshold = 1.5
lower_bound = Q1 - threshold * IQR
upper_bound = Q3 + threshold * IQR

# Remove outliers
data_no_outliers = data[(data["fare_amount"] >= lower_bound) &
    ↪(data["fare_amount"] <= upper_bound)]

# Visualize the 'fare_amount' distribution without outliers
sns.boxplot(x=data_no_outliers["fare_amount"])
plt.show()

```



```

[4]: data.plot(kind="box",subplots=True, layout=(7, 2), figsize=(15, 20))


```

```

[4]: Unnamed: 0      AxesSubplot(0.125,0.787927;0.352273x0.0920732)
fare_amount      AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_longitude AxesSubplot(0.125,0.677439;0.352273x0.0920732)
pickup_latitude  AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_longitude AxesSubplot(0.125,0.566951;0.352273x0.0920732)
dropoff_latitude  AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
passenger_count   AxesSubplot(0.125,0.456463;0.352273x0.0920732)

```

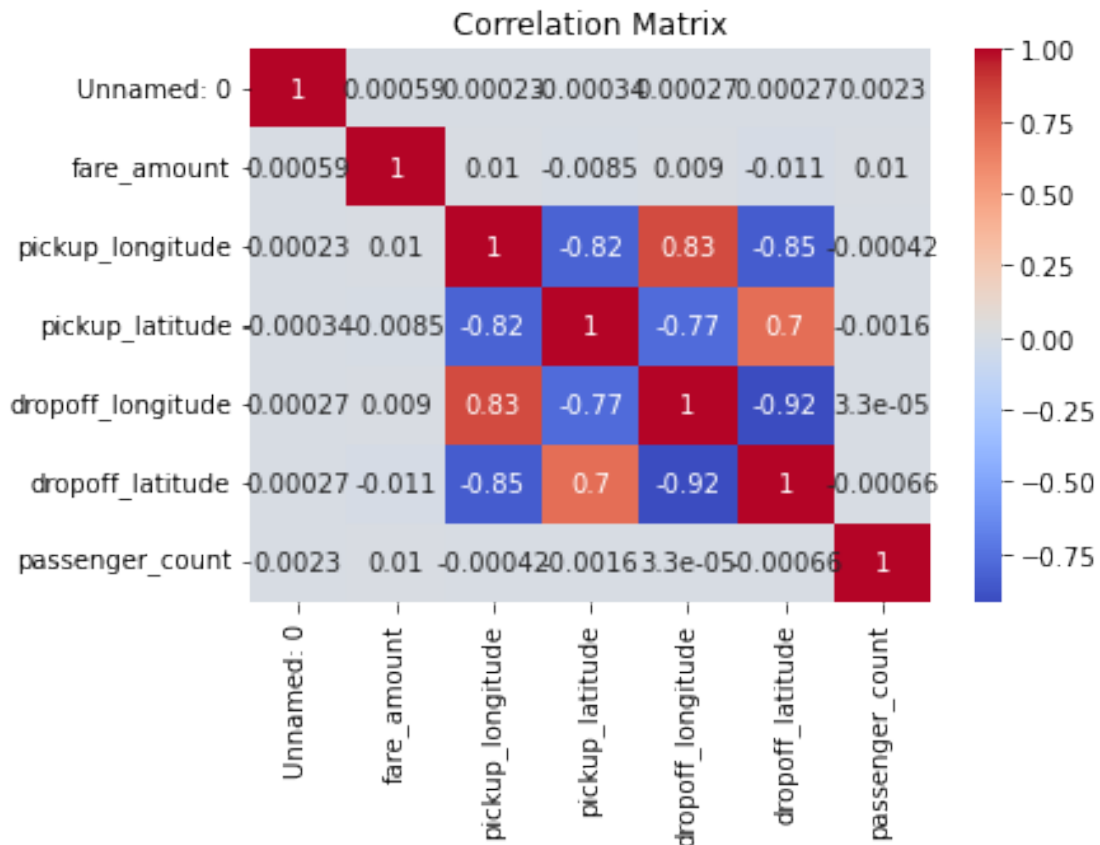
dtype: object



```
[5]: # 3. Check the correlation
# Determine the correlation between features and the target variable_
↳ (fare_amount).
numeric_data = data.select_dtypes(include=[int, float])

# Calculate the correlation matrix
correlation_matrix = numeric_data.corr()

# Plot the heatmap
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



```
[6]: # 4. Implement linear regression and random forest regression models
# Split the data into features and target variable
X = data[['pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
          'dropoff_latitude', 'passenger_count']]
y = data['fare_amount'] #Target

y
```

```
[6]: 0      7.5
     1      7.7
     2     12.9
     3      5.3
     4     16.0
     ...
    199995     3.0
    199996     7.5
    199997    30.9
    199998    14.5
    199999    14.1
    Name: fare_amount, Length: 199999, dtype: float64
```

```
[7]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
[8]: # Create and train the linear regression model
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
```

```
[8]: LinearRegression()
```

```
[9]: # Create and train the random forest regression model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
```

```
[9]: RandomForestRegressor(random_state=42)
```

```
[10]: # 5. Evaluate the models
# Predict the values
y_pred_lr = lr_model.predict(X_test)
y_pred_lr
print("Linear Model:", y_pred_lr)
y_pred_rf = rf_model.predict(X_test)
print("Random Forest Model:", y_pred_rf)
```

```
Linear Model: [11.29237916 11.29171388 11.5718662 ... 11.29183291 11.43252639
11.29190248]
```

```
Random Forest Model: [ 9.26    5.043 12.577 ... 6.8057 11.264  8.304 ]
```

```
[11]: # Calculate R-squared (R2) and Root Mean Squared Error (RMSE) for both models
r2_lr = r2_score(y_test, y_pred_lr)
rmse_lr = np.sqrt(mean_squared_error(y_test, y_pred_lr))
```

```
[12]: # Compare the scores
print("Linear Regression - R2:", r2_lr)
print("Linear Regression - RMSE:", rmse_lr)
```

```
Linear Regression - R2: 0.00034152697863043535
```

```
Linear Regression - RMSE: 10.197470623964248
```

```
[13]: r2_rf = r2_score(y_test, y_pred_rf)
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))

print("Random Forest Regression R2:", r2_rf)
print("Random Forest Regression RMSE:", rmse_rf)
```

```
Random Forest Regression R2: 0.701102567545169
```

```
Random Forest Regression RMSE: 5.576063738722296
```

```
[14]: # Overall Analysis

# The Random Forest Regression model has significantly improved the predictive
↳ performance.
# An R-squared (R2) value of approximately 0.701 and a Root Mean Squared Error
↳ (RMSE)
# of approximately 5.575 indicate that the Random Forest model is capturing a
↳ substantial portion
# of the variance in the target variable and providing more accurate
↳ predictions compared to the linear regression model.
```

```
[ ]:
```